

# ANSH KAUSHIK

Chicago, IL | 773-886-8967 | [ansh.kaushik1909@gmail.com](mailto:ansh.kaushik1909@gmail.com) | [www.linkedin.com/in/anshkaushik](https://www.linkedin.com/in/anshkaushik)

## SUMMARY

Computer Science graduate specializing in AI with Python development experience, seeking to apply academic ML implementations and workflow optimization skills in collaborative technical teams. Motivated to contribute hands-on expertise while expanding enterprise-scale system exposure

## EDUCATION

Illinois Institute of Technology, Chicago, IL

Aug 2024 - May 2026

**Master of Computer Science with Specialization in AI**, GPA: 3.67

*Machine Learning, Deep Learning, Natural Language Processing, Advanced Artificial Intelligence, Design and Analysis of Algorithms, Advanced Database Organization, Computer Networks, Big Data, Data Preparation and Analysis, Enterprise Web Applications*

Guru Gobind Singh Indraprastha University, Delhi, India

Aug 2019 - Jul 2023

**Bachelor of Technology in Computer Science and Engineering**, GPA: 3.81

## PROFESSIONAL EXPERIENCE

**Associate Software Engineer (AI/ML)**

Accenture, Gurugram, India

Sep 2023 - Jul 2024

- Prototyped a Python-based chatbot utilizing Azure OpenAI (GPT-3/3.5) and LangChain, cutting down employee search time for internal documentation by 25% and streamlining workflow automation for 500+ users
- Implemented generative AI-powered text summarization feature in a Power BI dashboard leveraging GPT API, decreasing report preparation turnaround by 30 minutes per analytics request
- Acknowledged by management for developing clear technical documentations supporting project adoption by non-technical teams

**Data Science Intern**

Alphec Global Private Limited, Noida, India

Jan 2023 - Jun 2023

- Cleaned and normalized structured datasets using Python (Pandas), resolving 500+ missing values and outliers to prepare data for predictive modeling tasks
- Deployed logistic regression models for customer behavior analysis leveraging scikit-learn, accomplishing 72% accuracy on a stratified test set (12K records) after addressing class imbalance through class weights
- Optimized Snowflake SQL queries, decreasing dataset processing time by 30% to accelerate model iteration cycles

**Data Analysis IT Intern**

Lumax Industries Ltd, Gurugram, India

Jul 2022 - Sep 2022

- Collaborated within cross-functional Agile teams and volunteered to resolve 15+ system incidents through root-cause analysis, revamping workflow logic stability
- Shadowed senior developers on demand forecasting modules, observing time-series analysis techniques applied to supply chain planning and development of knowledge graphs and clustering techniques for supply chain optimization

## PROJECTS

**Conversational Text-to-SQL Query And Visualization Assistant**

- Created an AI-enabled Streamlit app, turning complex natural language queries into SQL and visualizations with <10% error rate, supporting real-time interactive analysis to data across 7 interconnected datasets
- Reduced average data exploration and visualization time by 60%, simplifying SQL query execution speed by 30% through prompt optimization and validation techniques, resulting in more efficient data retrieval
- Implemented secure multi-user cloud deployment with session management and query favorites, boosting analysis efficiency by 35%

**Cloud-Based Banking Customer Classification System**

- Developed and fine-tuned classification models (Logistic Regression, SVM, Random Forest) on merged banking datasets (~10,000 records), achieving 82% accuracy in predicting likely loan customers
- Streamlined workflow by automating data preprocessing and model selection, cutting manual effort in half, and revising model precision and recall by 15%
- Built and containerized a Flask API with Docker for best model, deployed on AWS EKS with CI/CD to enable scalable, automated live predictions in a simulated production environment

### Multi-Class Text Classification With Transformer Models

- Fine-tuned ALBERT and DistilBERT transformer models on the DBpedia ontology dataset (560K training samples), achieving up to 86% classification accuracy across 14 categories and diminishing inference latency by 25%
- Performed extensive data preprocessing and visualization using Python libraries (pandas, matplotlib, numpy) to analyze text features, improving model input quality, increasing validation accuracy by 5%
- Conducted comparative performance analysis with baseline BERT models on DBpedia and Emotion datasets, demonstrating up to 18x parameter reduction and 50% smaller model size deploying knowledge distillation techniques

### Hospital Resource Optimization Via LOS Prediction

- Developed a predictive pipeline using Snowflake and AWS SageMaker, processing 50k+ patient admission records to estimate patient length of stay, improving prediction accuracy by 15% over baseline models
- Conducted comparative analysis of regression models (Linear, Random Forest, XGBoost), with feature engineered XGBoost outperforming others by 15% lower MAPE, reducing overfitting and enhancing real-world applicability
- Automated data ingestion, prediction storage, and scheduling with Snowflake connectors and AWS job orchestration, reducing manual workflows by 90%

## SKILLS

---

- **Programming Languages:** Python, C/C++, R, TypeScript, JavaScript
- **Frameworks & ML Libraries:** PyTorch, TensorFlow, Keras, scikit-learn, Transformers, pandas, numpy, OpenCV, LangChain, Streamlit, Flask, React
- **Cloud and Databases:** Amazon Web Services (AWS), Azure, Databricks, Snowflake, MySQL, MongoDB, Firebase, Hadoop, Kafka, Vector Databases
- **DevOps & Tools:** Git, GitHub, Docker, Kubernetes, CI/CD, Jenkins, VSCode, Jupyter, Selenium, Agile

## CERTIFICATIONS

---

- Machine Learning Specialization - Stanford Online / DeepLearning.AI
- Deep Learning Specialization - DeepLearning.AI (In progress, expected completion: December 2025)